Document discovery and inclusion: methods

How the second collection round found, verified, de-duplicated, acquired and admitted governance instruments, and how the instruments are counted

Centre for Enterprise Innovation (CEI) AI Robotics project, with Claude Code

31 October 2026

- 1. Rationale and design
- 2. Sources
- 3. Candidate schema and verification
- 4. Merging and de-duplication
- 5. Text acquisition
- 6. Annotation, inclusion and counting
- 7. Quality checks
- 8. Yield by source
- 9. Limitations of the discovery method

1. Rationale and design

The first collection was assembled by eight regional teams in two exercises: a *sprint* for robot-specific sources and a *database* of general AI-governance instruments. Its coverage audit showed the limitation of that design for the question the study asks. Re-computed against the current reference list of 39 instruments known to govern robots, drones or automated vehicles, the first collection held 11 of them as principal instruments with a content-analysed capture (12 on a looser title match), and it held 65 texts of instrument class in the embodied-AI set, 5 of them legally binding. The instruments that bind robots in practice (aviation rules, road-traffic law, machinery and product-safety law, workplace-safety rules, technical standards) are catalogued in legal databases that topical web search rarely surfaces, and they are spread across many jurisdictions and languages.

The second collection round therefore inverted the search strategy. Instead of asking regional teams to search the web for material *about* embodied AI, it (a) queried official legal databases and catalogues programmatically for instruments whose text or subject tags name the technologies, and (b) briefed nine specialised collector agents to work from legal repositories, treaty archives, standards catalogues and regulator sites with explicit lists of the instruments and instrument families to look for. The collector agents were AI agents (Claude subagents running in one Claude Code session on 30 October 2026, documented in `collection_round2/ROUND2_RUNBOOK.md`); no person searched for, read or labelled any document in either round. Discovery was separated from judgement: every candidate that could be fetched was annotated by the same metadata-blind procedure as the rest of the collection, and only the annotation decides whether a document is an instrument, whether it concerns embodied AI and how it is classified. A set of documented title and capture rules then assigns each text of instrument class a counting unit (section 6.3), so that the headline counts refer to principal instruments rather than to documents. Candidates that could not be fetched are recorded in an inventory and are not counted anywhere.

2. Sources

2.1 Scripted harvests of legal databases

Seven databases were harvested by script (eight runs; `collection_round2/scripts/harvest_apis.py`). Each harvester returns rows in the common candidate schema (section 3). The metadata fields of harvested rows (instrument type, binding nature, status) are constants set per source in the harvester, not judgements about the individual document.

- **EUR-Lex through the CELLAR SPARQL endpoint**, in two ways: (i) a title search over the English expression titles of regulations, directives, decisions (including implementing and delegated acts), resolutions, proposals, communications and recommendations, using 25 terms covering robots, unmanned aircraft and drones, automated and autonomous driving, driver assistance, machinery, toy safety, product liability, medical devices, lifts, cyber resilience, general product safety, U-space, type-approval, autonomous and lethal autonomous weapons, and artificial intelligence; (ii) a subject search for every work tagged with the EuroVoc concepts *robotics*, *industrial robot*, *autonomous vehicle*, *drone* and *artificial intelligence*, which is independent of title wording. Corrigenda, case law, national transposition measures, parliamentary questions and administrative notices (merger notifications, state aid, vacancies) were excluded.
- **legislation.gov.uk**, title search over 15 terms (automated vehicles, unmanned aircraft, drone, robot, autonomous, air navigation, product safety, machinery, vehicle technology, self-driving, remotely piloted, toy safety, medical devices), excluding retained EU instruments.
- **Japan's e-Gov law API**, full-text search with nine Japanese terms (specified automated operation, unmanned aircraft, remotely operated small vehicles, industrial robots, automated driving, robots, automated operation devices, care robots, drones), returning the consolidated law text.
- **China's State Council policy library** (gov.cn), full-text search with twelve Chinese terms (humanoid robots, embodied intelligence, robots, automated driving, unmanned aircraft, intelligent connected vehicles, drones, low-altitude economy, elderly-care robots, surgical robots, industrial robots, service robots).
- **Australia's Federal Register of Legislation** (OData API), title search with nine terms.
- **The US Code of Federal Regulations** (eCFR search API), quoted-phrase searches for twenty terms aggregated to the level of a CFR part, retaining parts in titles where such rules plausibly live or with several matching sections, and retrieving the part text through the versioner API.
- **Brazil's Câmara dos Deputados open-data API**, keyword search in Portuguese (low yield; the Latin America collector covered Brazil from the Senate and Chamber sites instead).

Harvests of the Swiss Fedlex gazette and the Korean law portal were attempted and dropped (JavaScript-only pages and a registration-only API respectively). Several further databases could not be reached from the execution environment (GovInfo and regulations.gov without an API key, the Federal Register API, FAOLEX, NATLEX, OECD STIP Compass, New Zealand legislation, Taiwan and Korean English law portals).

2.2 Collector agents

Nine collector agents ran in parallel, each with a written brief that stated the scope, the target number of rows, the candidate schema, the list of round-1 holdings to exclude, a long enumeration of instruments and instrument families to look for in that scope, the databases and hubs to use, and the rule that every primary URL must have been verified to resolve. The scopes were: the United States (federal, state and city); the European Union, the United Kingdom, Ireland and the Nordic countries; continental European national instruments; East Asia including Chinese provinces and municipalities; South and Southeast Asia and Oceania; the Middle East, Africa, Latin America and Canada; technical standards (ISO/IEC, IEEE, SAE, UL, ASTM, BSI and national bodies); intergovernmental, military and national-strategy instruments (the CCW process on lethal autonomous weapons, UN resolutions and reports, NATO, REAIM, national defence AI policies, national robotics strategies); and medical, workplace, consumer and care-robot rules (medical-device guidance, occupational-safety rules on industrial and collaborative robots, toy and companion-robot rules, municipal delivery-robot rules). A tenth agent collected the public scope or abstract text of standards whose catalogue pages block automated access.

Agents used web search for discovery while a shared search budget lasted (it was exhausted part-way through every brief) and thereafter worked from official databases and site listings (for example the National People's Congress law database, the SAMR standards platform, law.go.kr, the Sejm and Legilux legal-data interfaces, the UK Parliament committees API, the UNODA document library, Senate and Chamber open-data APIs). Each agent returned a schema-exact CSV written in batches and a report listing what it had looked for and could not find; those reports are reproduced in condensed form in the package (`collection_round2/COLLECTOR_REPORTS.md`) because they bound what the round can claim.

3. Candidate schema and verification

Every candidate row carries 17 fields: identifier; original and English titles; jurisdiction (ISO 3166-1 alpha-2, EU or INT) and subnational unit; issuer; year; instrument type (13 values from statute to court decision); binding nature (legally binding, voluntary soft law, voluntary or mandated technical standard, not applicable); lifecycle status (eight values); robot types addressed (from the study's 23-category taxonomy); language; primary and secondary URLs; the source database; a confidence level (high if verified at the primary URL, medium if verified through a secondary source, low if inferred); and a one-line note. The inclusion rule was governance instruments only (laws, regulations, bills, standards, strategies, official guidelines, resolutions, formal national positions, official reports), one row per instrument, preferring instruments with substantive provisions on physical AI. Collectors were asked to include medium- and low-confidence rows with their official URL rather than omit them. Across the merged inventory, 977 rows were recorded as high confidence, 1,008 as medium and 90 as low. The collector metadata were recorded by the agents from catalogue or gazette pages; they were never shown to the annotation model and are used only for the inventory and for the consistency check in section 7.1.

4. Merging and de-duplication

Candidate files were merged with a fixed procedure (`merge_candidates.py`). Rows without a title or a primary URL were dropped, as were the Fedlex gazette hits and EUR-Lex corrigenda.

For the title-searched harvests only (EUR-Lex by title, legislation.gov.uk, Fedlex, Australia), a keyword relevance filter kept titles naming the technologies or the relevant product-safety regimes and removed administrative noise (anti-dumping, customs, mergers, state aid, appointments, autonomous-region place names); full-text harvests and collector rows were not filtered by title, so that the annotation, not a keyword list, judges them. Titles were normalised (Unicode NFKC, case, parentheticals, punctuation, leading articles) and URLs canonicalised (scheme, `www`, fragment, trailing slash). Duplicates were removed within the round on URL and then on normalised title plus jurisdiction, and against the round-1 collection on URL or on normalised title of at least 20 characters. Each surviving candidate received a stable identifier derived from its canonical URL, so that re-runs added only new rows. The merged inventory holds 2,075 candidates: 1,157 from collector agents and 918 from scripted harvests, across 88 jurisdiction codes.

Two further de-duplication passes run when the round-2 documents are assembled with the round-1 documents into the master table: documents with an identical normalised title in the same jurisdiction and years within one are merged (the fullest capture is kept; 119 rows across the whole collection, 78 of them from round 2), and a round-2 document is removed when a round-1 document has a token-sort title similarity of at least 90 with comparable length, or when both titles open with the same EU act number and neither is an amendment, opinion, guideline or report about the act (1 round-2 row). A looser similarity rule tried first would have removed 69 distinct documents by matching them to short generic round-1 titles; every removal is listed in `data/duplicates_removed.csv`. In all, 79 of the 1,720 round-2 texts were removed in assembly and 1,641 entered the master table. Two further rules operate on the assembled table at the level of the counting unit rather than the document (section 6.3): texts of instrument class whose files are byte-identical are one capture (`duplicate_text`), and texts whose titles differ only in year, act number, version or amendment wording within one jurisdiction are grouped into an *instrument family* (`family_key`), so that yearly renewals and successive versions can be counted once where that unit is wanted.

5. Text acquisition

Texts were fetched from the primary URL with the secondary URL as fallback (`fetch_texts.py`), with routing rules for sources that block or mis-serve automated clients: EUR-Lex links are resolved through the Publications Office's CELLAR service (XHTML, falling back to PDF or HTML for older acts); UN and UNODA links are resolved through the UN document API by document symbol; legislation.gov.uk items are read as XML; ISO catalogue pages are read from the ISO committee mirror; hosts that reject the execution environment are skipped by exact host name. PDFs were converted with pdftotext; HTML with a content-area extractor; Japanese and British XML by tag stripping. A text was accepted when it had at least 80 words or 400 CJK characters. Where a standards catalogue page could not be fetched, the scope or abstract text collected by the standards agent was used instead and flagged in the registry.

Text was obtained for 1,720 candidates (1,704 fetched, 16 standards scope texts). The 355 candidates without text (272 HTTP errors or script-only pages, 46 on hosts that block the environment, 37 failed transfers) remain in the inventory with their metadata. After the review of the previous version, every failed fetch was retried with the improved routes; this recovered-fetch pass added 4 texts (a Tanzanian amendment to the civil-aviation rules on remotely piloted aircraft, the catalogue page of the Chinese geofencing standard GB/T 43370-2023, and two Korean ministry announcement pages on autonomous ships and

mobility). The hosts accounting for most losses were eur-lex.europa.eu for items with no English expression (59 candidates), the National People's Congress law database (27), the UNODA library for documents without a UN symbol (17), Légifrance (17), the Korean law portal (17) and the Taiwanese law portal (11). Of the 355 candidates without text, 189 were recorded by their collectors as legally binding; they include the instruments named in section 9 as the first targets of a third round.

6. Annotation, inclusion and counting

6.1 Annotation

Every fetched text was annotated with the procedure used for the governance documents of the first collection: the 39-field codebook, a fixed JSON schema, Claude Sonnet 5.5 through the Claude Code command-line interface, the collector's cover block stripped and only minimal metadata shown, and an extended excerpt of up to 110,000 characters (the first 100,000, 6,000 from the middle, 4,000 from the end). The model, prompt hash, excerpt policy and cost are stored with each annotation, and the exact model input for every round-2 document is in `annotations/round2/_inputs`; `data/run_manifest.json` lists models, prompt hashes, policies, counts and cost per annotation folder (1,736 round-2 annotation calls at about 126 US dollars, including the re-annotations described next).

The excerpt cleaner used for the previous version removed short lines near the top of a text that carried no digit, capital letter or sentence punctuation, a rule written for navigation menus in Latin-script web pages. Chinese, Japanese and Korean paragraphs have no whitespace-delimited words and no ASCII full stop, so the rule deleted most of their text before the model saw it; the reviewer who found the fault measured that a median of 24% of each Chinese document reached the model. The cleaner was made character-aware (lines that carry CJK characters or end in CJK punctuation are never dropped; policy tag `cjk_clean=v4`), and every Chinese, Japanese and Korean document was re-annotated: 376 round-2 and 46 round-1 documents. Of these 422 documents, 57 changed embodied-set membership, and the number rated as not about AI fell from 205 to 160. Sixteen standards inputs that carried the collecting agent's header were re-annotated with the header stripped.

6.2 Inclusion

Inclusion follows the study's standing definitions. A document is *usable* if its text was obtained and the annotation did not judge the capture to be an error page or otherwise unusable (1,626 of the 1,641 round-2 documents in the master table). It belongs to the *embodied-AI set* if it is rated at least moderately relevant to embodied-AI governance or carries a robot-type label, and is not judged to be unrelated to AI (928 round-2 documents). It is a *text of instrument class* if its document class is binding instrument, soft-law instrument or proposal/draft (1,351 round-2 documents; 745 in the embodied set), and *binding* if its class is binding instrument (286 round-2 texts in the embodied set). Documents rated as not about AI at all, mostly harvested legal texts in which the technology terms appear incidentally, fall outside both analysis sets.

Four rules added after the review of the previous version correct status and class labels that the annotation got wrong at scale: standards still in development (AWI to FDIS stages) are not published instruments and are classed as drafts under development; mandatory Chinese GB standards (not GB/T or GB/Z) are binding; CCW working papers, chair's summaries and rolling

texts are position and process documents, not draft instruments; and a proposal whose adopted act is in the collection is superseded. In addition, 157 documented document-level corrections (on 91 documents, 156 of the rows concerning round-2 documents) were converted from the governance expert's audit of the previous version and are applied last (`data/manual_overrides_v4.csv`: counting unit 50 rows, relevance 32, robot categories 32, status 17, document class 14, year 7, binding nature 5). That audit inspected a risk-stratified sample of 159 rows from the previous version's instrument tables and found at least one error in 108 of them (68%) and a material error (duplicate or dependent act, not an instrument, wrong class, wrong status, or not about embodied-AI governance) in 98 (62%); because the sample was risk-stratified, these are upper bounds for the tables as a whole. The audit table released with the package holds 165 rows, 119 of them with an error code.

6.3 Counting unit

The headline unit is the *principal instrument*. Among texts of instrument class, documented title and capture rules (`scripts/07c_v4_units_and_tiers.py`) set aside commencement orders, levy and fee instruments, individual exemptions, directions, approvals and permits, definitions-only parts and tables of contents (`ancillary`); EEA Joint Committee incorporation decisions (`incorporation`); language-version corrigenda (`language_copy`); byte-identical duplicate texts (`duplicate_text`); catalogue, register and landing pages (`catalogue_page`); metadata stubs under 150 words whose capture is not a substantive document (`stub`); and dependent acts identified by the audit (`dependent`). Amendments count as instruments. In the usable round-2 documents, the 745 embodied-set texts of instrument class resolve into 403 principal texts and 41 amendments (444 principal instruments), 218 catalogue pages, 35 stubs, 30 ancillary instruments, 10 dependent acts, 6 incorporation decisions and 2 language copies; the 286 binding embodied-set texts resolve into 190 principal binding instruments, 164 of them rated robot-relevant (relevance 2 or 3), 135 with autonomy or AI provisions in the text and 10 naming AI explicitly. Across the whole collection, table T53b gives the same counts under successively stricter units: 810 embodied-set texts of instrument class, 505 principal instruments, 408 rated robot-relevant, 395 whose text contains autonomy or AI provisions and 153 naming AI explicitly; for binding instruments, 291 texts, 195 principal instruments, 168 robot-relevant, 194 when counted once per instrument family, 140 with autonomy or AI provisions and 15 naming AI explicitly. Every count in the study names its unit when it is not the principal unit; comparisons with the first collection are made on the comparable unit (all texts of instrument class, table T53) and are read as a consequence of collection design, not as a change in the world.

The construct follows from these tiers. The embodied-AI set is defined by annotation relevance or a robot-type label, and most binding instruments that reach robots are machine-class rules (aviation, road vehicle, machinery, product safety) with no AI content: of the 195 principal binding instruments in the embodied set, 140 contain autonomy or AI provisions and 15 name AI explicitly. The object of study is therefore the governance of the machine classes in which embodied AI is deployed, with the AI-content tiers reported alongside every headline count.

6.4 Capture tier and dating

Statistics about what instruments specify in depth (the principle-depth and mechanism tables, the autonomy-level and physical-safety shares) use content-analysed captures with at least 150 words; a capture is content-analysed when the annotation judged it a substantive document or a partial excerpt. Standards scope pages are reported as a separate

measurement and never pooled with substantive texts: 200 of the 260 round-2 technical standards are catalogue or landing pages, and the depth a scope page can show is the depth of a scope statement. Consolidated statutes harvested from legal databases carry the enactment year or the version year as the annotation found it; the time analyses are restricted to 2014 to 2026, treat 2026 as partial, and do not report “first appearance” years for sector law.

7. Quality checks

7.1 Consistency between collector metadata and annotation

The collectors’ metadata were compared with the blind annotation of the text for the usable round-2 documents: agreement was 92.9% on whether an instrument is legally binding (kappa 0.85, n = 1,608), 98.1% on jurisdiction (n = 1,624), 93.8% on year within one (n = 1,542) and 76.7% on document class (kappa 0.66, n = 1,608), the main disagreements being guidelines that the annotation reads as binding and strategies or consultations that it reads as communications. This is a consistency check, not a validation: the collectors were AI agents of the same model family as the annotator, the harvested rows carry one constant label per source, and no person recorded any of the reference values.

7.2 Round-2 reliability

A stratified sample of 160 round-2 documents was re-annotated by a second model (Claude Opus 5.5) and by a same-model re-run under the same policy. The median kappa over single-label and ordinal fields was 0.86 against the second model and 0.91 for the re-run, next to 0.85 and 0.91 for the round-1 samples; the weakest round-2 fields against the second model were page quality (0.53), stance toward regulation (0.65), publication stage (0.66), date basis (0.70) and instrument status (0.75). Unweighted and level-2 variants are in table S09b. No human validation exists: no person read or labelled any document.

7.3 Coverage audit

The reference list now has 39 entries (UN Regulations 157 and 171 as separate rows, the Taiwanese industrial-robot standards dated 2018, and a stricter title pattern for US state automated-vehicle statutes). Under the strict criterion, an instrument is captured when a principal instrument with a content-analysed capture matches its title pattern: 22 of 39 are captured (11 for the first collection). Under the looser title-family match used previously, 31 of 39 are captured (12 for the first collection). The 17 still missing under the strict criterion are the German Autonomous Driving Act of 2021, the Korean Intelligent Robots Development and Distribution Promotion Act, the revised EU Product Liability Directive (EU) 2024/2853, the Shenzhen intelligent-connected-vehicle regulations, the Beijing humanoid-robot and embodied-intelligence measures, ISO/TS 15066, IEC 61508, ISO 26262, ISO 21448, ISO/IEC 42001, IEEE 7001, ANSI/RIA R15.06, ANSI/A3 R15.08, the Korean AI Basic Act, the OECD AI Principles, the Taiwanese industrial-robot hazard-prevention standards and UN Regulation No. 171 (table T25).

7.4 Precision audit

An independent audit re-read a simple random sample of 100 of the 299 documents that the previous version counted as binding embodied-AI instruments and judged each against its

own text (`review/v3_reviews/precision_audit.md`). Of the 100, 67 were distinct instruments (95% Wilson interval 57 to 75), 97 were binding, 56 contained substantive embodied provisions (46 to 65), 67 contained autonomy or AI content and 14 named AI explicitly (9 to 22). All three of distinct, binding and substantive held for 54 (44 to 63), which implies about 161 such instruments (132 to 190) in that population of 299 rather than 299. On the same 100 documents the rule-based unit of this version passes 61 (51 to 70), agreeing with the audit on 77 documents (kappa 0.53; table T62). The audit measures precision only; it does not consolidate instrument families and it does not add instruments the collection missed, so the precision and coverage corrections pull in opposite directions and are reported separately. The precision of every instrument count in this study should be read as an upper bound.

7.5 Numbers audit

An independent numbers audit recomputed every quantitative statement in the narrative documents from the master table and the tables, and the discrepancies it found were corrected and logged (`review/v4_numbers_audit*.md`). The numbers auditor, like the collectors and the six reviewers of the previous version, was an AI agent.

8. Yield by source

Source	Candidates	Texts	Usable	Principal embodied instruments
EUR-Lex, title search	265	251	251	33
EUR-Lex, EuroVoc concepts	256	224	222	30
Technical standards collector	244	231	231	58
China State Council policy library	192	188	188	41
United States collector	176	131	85	50
Intergovernmental, military and strategy collector	148	115	79	29
East Asia collector	146	76	75	20
EU, UK, Ireland and Nordics collector	102	77	75	37
Middle East, Africa, Latin America, Canada collector	97	69	67	37
South and Southeast Asia and Oceania collector	95	64	63	22
Continental European collector	85	46	45	25

Source	Candidates	Texts	Usable	Principal embodied instruments
Japan e-Gov law API	64	64	64	22
Medical, workplace, consumer and care collector	64	53	52	20
US eCFR	55	55	53	12
legislation.gov.uk	46	37	37	7
Australian Federal Register of Legislation	38	38	38	0
Brazil Câmara API	2	1	1	1
All sources	2,075	1,720	1,626	444

Counts are for usable round-2 documents on the principal-instrument unit (table T50 gives the full breakdown, including embodied-set documents and all principal instruments per source). Yield differs by design: the standards collector and the regional collectors yield embodied instruments for a large share of their candidates, while the broad EUR-Lex and gov.cn harvests contribute many general-AI or non-AI legal texts that the annotation sets aside. The Australian register illustrates the effect of the counting unit: its 38 candidates are all usable and 32 are in the embodied set, but they are register pages and individual approvals, directions and exemptions, so none is a principal instrument. The two strategies remain complementary: scripted harvests supplied most of the binding EU, Japanese, Chinese and US federal rules, and the collectors supplied national and subnational instruments, standards and treaty-body material that no single database indexes.

9. Limitations of the discovery method

Coverage is bounded by what the execution environment could reach. Portals that block automated access or serve script-only pages (among them the Korean and Taiwanese law portals, the National People's Congress law database, Hong Kong e-Legislation, Légifrance, defense.gov, the Indian civil-aviation host, the Singapore statutes site, the ISO store, UNECE, ICAO and OSHA) could be cited but not read, and 355 verified candidates therefore remain outside every count, 189 of them recorded by their collectors as legally binding. Instruments identified and verified by the collectors but still unread include the Hangzhou regulations on the embodied-intelligent-robot industry (2026); the Shenzhen regulations on intelligent connected vehicles (2022), on the AI industry (2022) and on the low-altitude economy (2024); the Hong Kong Road Traffic (Autonomous Vehicles) Regulation (2024); the Korean Intelligent Robots Development and Distribution Promotion Act and its enforcement decree, and the enforcement decree and rule of the Korean autonomous-vehicles commercialisation act (2020); US DoD Directive 3000.09 on autonomy in weapon systems (2023) and California SB 243 on companion chatbots (2025); India's Drone Rules 2021; the EU Product Liability Directive (EU) 2024/2853, the Toy Safety Regulation (EU) 2025/2509 and the Toy Safety Directive 2009/48/EC; the French ordinance 2021-443 and decree 2021-873 on criminal liability for automated vehicles; the Swiss DETEC ordinance on aircraft of special categories (2022); and the standards ANSI/A3 R15.06-2025 and GB/T 48317.1-2026 on humanoid robots. These are the first targets of a third collection round; their absence means that at least one

robot-titled statute, the horizontal consumer-product regimes most relevant to companion robots, and the main Korean instruments are missing from every count in this study. Because availability depends on which jurisdictions expose their law through reachable databases, the collection over-represents the EU, the United Kingdom, Japan, China's central government, Australia and the US federal level relative to jurisdictions whose instruments are not so indexed.

The shared web-search budget was exhausted in every collector brief, so discovery leaned on databases with usable interfaces. Title-searched harvests depend on the keyword lists given above and miss instruments whose titles do not name the technology; full-text harvests depend on the same vocabulary in the body. Technical standards were annotated on their public scope text rather than the normative text, and scope pages are kept apart from substantive captures in every depth statistic; the copyright status of the mirrored scope texts limits what the package can redistribute. The collectors' metadata were used only for the consistency check and the inventory, never for classification, which rests on the model annotation with the reliability documented in section 7.2 and the supplement; the collectors, the reviewers of the previous version, the precision auditor and the numbers auditor were all AI agents, and no person read or labelled any document. Instrument counts are counts of captured texts under documented rules: the precision audit shows that about half of the documents counted as binding embodied instruments in the previous version were distinct, binding and substantive, and the rule-based unit agrees with that audit on 77 of 100 documents, so every count remains an upper bound on precision and, because of the unread candidates, a lower bound on coverage within this collection. Statutes harvested as consolidated texts carry the year the annotation found, which may be the enactment year or the consolidation year. Three criticisms of the previous version were not addressed: no human validation of any label exists (no human coders were available in this setting); statutes longer than the 110,000-character excerpt were not re-annotated in full, so provisions beyond the excerpt may be coded from the model's prior knowledge of the instrument; and stricter title merging of near-duplicate pairs was not applied, because distinct parts of one regime are legitimately distinct documents, with instrument families offered as the alternative unit.